




c

Method	Head score	Prunes in training*	Gradual fade	Reopens	exact count never broke
Hemodynamic attenuation (ours)	$\min_W L_{-h} - \min_W L$ Optimal Brain Surgeon (Hassibi & Stork 1993)	●	●	●	29/30 30/30
<i>Published methods</i>					
Kurtic et al. 2023 (ZipLM)	$\min_W L_{-h} - \min_W L$ Optimal Brain Surgeon (Hassibi & Stork 1993)	—	—	—	24/30 0/30
Michel et al. 2019	$ \partial L / \partial g_h $ gradient of the head gate	—	—	—	28/30 0/30
Voita et al. 2019	$L + \lambda \sum_h P(g_h \neq 0)$ a learned gate per head (L0 penalty)	—	—	—	9/30 0/30
<i>Controls: our method with one part removed</i>					
no fade (instant closing)	$\min_W L_{-h} - \min_W L$ Optimal Brain Surgeon (Hassibi & Stork 1993)	●	—	●	30/30 0/30
no re-fit	$L_{-h}(W) - L(W)$ Optimal Brain Damage (LeCun et al. 1989)	●	●	●	2/30 30/30
no score (random choice)	random head	●	●	●	30/30 6/30

Tasks needing 2, 4 and 8 heads, 10 seeds each; every method run the same way here. *as published. ● yes — no

