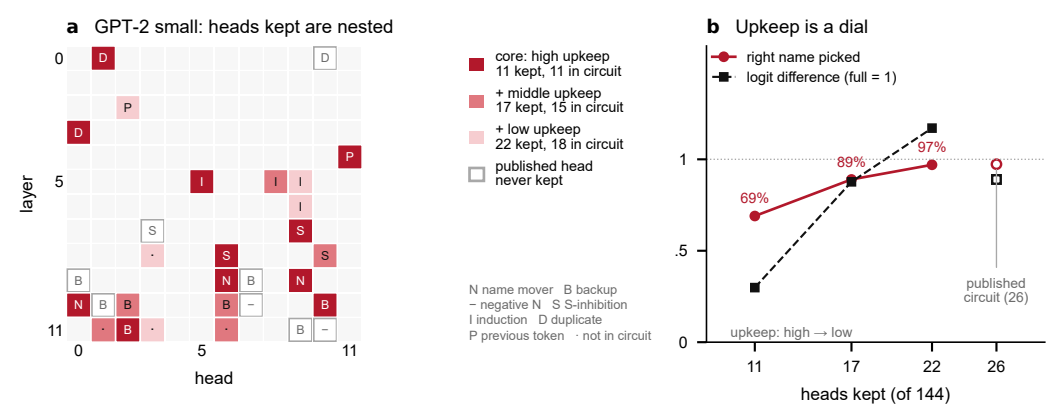
a
GPT-2 small: heads kept are nested
layer
head
0
5
11
core: high upkeep
11 kept, 11 in circuit
+ middle upkeep
17 kept, 15 in circuit
+ low upkeep
22 kept, 18 in circuit
published head
never kept
N name mover B backup
− negative N S S-inhibition
I induction D duplicate
P previous token · not in circuit
b
Upkeep is a dial
right name picked
logit difference (full = 1)
69%
89%
97%
published
circuit (26)
upkeep: high → low
heads kept (of 144)
1
.5
0
11
17
22
26